

Assignment no. 04

Aim-

1. Linear Regression : Univariate and Multivariate
2. Least Square Method for Linear Regression
3. Measuring Performance of Linear Regression
4. Example of Linear Regression
5. Training data set and Testing data set

```
In [9]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [23]: x=np.array([95,85,85,70,60])
y=np.array([85,90,70,64,70])
model= np.polyfit(x, y, 1)
model
```

```
Out[23]: array([ 0.53766234, 33.32467532])
```

```
In [24]: predict = np.poly1d(model)
predict(65)
```

```
Out[24]: 68.27272727272727
```

```
In [25]: y_pred= predict(x)
y_pred
```

```
Out[25]: array([84.4025974 , 79.02597403, 79.02597403, 70.96103896, 65.58441558])
```


```
In [26]: from sklearn.metrics import r2_score
r2_score(y, y_pred)
```

```
Out[26]: 0.4516887333445776
```

```
[27]: y_line = model[1] + model[0]* x
plt.plot(x, y_line, c = 'r')
plt.scatter(x, y_pred)
plt.scatter(x,y,c='r')
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x1e75c510c90>
```

In



```
[28]: import ssl from sklearn.datasets import fetch_california_housing
ssl._create_default_https_context = ssl._create_unverified_context
california = fetch_california_housing(download_if_missing=True) X
= california.data y = california.target california
```

```
Out[28]: {'data': array([[ 8.3252, 41., 6.98412698, ..., 2.55
555556,
37.88, -122.23 ],
[ 8.3014, 21., 6.23813708, ..., 2.10984183,
37.86, -122.22 ],
[ 7.2574, 52., 8.28813559, ..., 2.80225989,
37.85, -122.24 ],
...,
[ 1.7, 17., 5.20554273, ..., 2.3256351,
39.43, -121.22 ],
[ 1.8672, 18., 5.32951289, ..., 2.12320917,
39.43, -121.32 ],
[ 2.3886, 16., 5.25471698, ..., 2.61698113,
39.37, -121.24 ]]),
'target': array([4.526, 3.585, 3.521, ..., 0.923, 0.847, 0.894]),
'frame': None,
'target_names': ['MedHouseVal'],
'feature_names': ['MedInc',
'HouseAge',
'AveRooms',
'AveBedrms',
'Population',
'AveOccup',
'Latitude',
'Longitude'],
```

In

```
'DESCR': '.. _california_housing_dataset:\n\nCalifornia Housing dataset\n-----\n\n**Data Set Characteristics:**\n\n: Number of Instances: 20640\n\n: Number of Attributes: 8 numeric, predictive attributes and the target\n\n: Attribute Information:\n\n- MedInc median income in block group\n\n- HouseAge median house age in block group\n\n- AveRooms average number of rooms per household\n\n- AveBedrms average number of bedrooms per household\n\n- Population block group population\n\n- AveOccup average number of household members\n\n- Latitude block group latitude\n\n- Longitude block group longitude\n\n: Missing Attribute Values: None\n\nThis dataset was obtained from the StatLib repository.\nhttps://www.dcc.fc.up.pt/~ltorgo/Regression/cal_housing.html\n\nThe target variable is the median house value for California districts, expressed in hundreds of thousands of dollars ($100,000).\n\nThis dataset was derived from the 1990 U.S. census, using one row per census block group. A block group is the smallest geographical unit for which the U.S. Census Bureau publishes sample data (a block group typically has a population of 600 to 3,000 people).\n\nA household is a group of people residing within a home. Since the average number of rooms and bedrooms in this dataset are provided per household, these columns may take surprisingly large values for block groups with few households and many empty houses, such as vacation resorts.\n\nIt can be downloaded/loaded using the\nfunc:`sklearn.datasets.fetch_california_housing` function.\n\n.. topic:: References\n\n- Pace, R. Kelley and Ronald Barry, Sparse Spatial Autoregressions, Statistics and Probability Letters, 33 (1997) 291-297\n'}
```

```
[29]: data = pd.DataFrame(california.data)
data.columns = california.feature_names
data.head()
```

Out[29]:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25

```
In [30]: data['PRICE'] = california.target
data.isnull().sum()
```

```
Out[30]: MedInc      0
HouseAge    0
AveRooms    0
AveBedrms   0
Population  0
AveOccup    0
Latitude    0
Longitude   0
PRICE       0
dtype: int64
```

In

```
In [31]: data.isnull().sum() from sklearn.model_selection
import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(X, y, test_size=0.2, random_
import sklearn
from sklearn.linear_model import LinearRegression
lm = LinearRegression() model=lm.fit(xtrain,
ytrain)
```

```
In [32]: ytrain_pred = lm.predict(xtrain)
ytest_pred = lm.predict(xtest)
df=pd.DataFrame(ytrain_pred,ytrain)
df=pd.DataFrame(ytest_pred,ytest)
from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(ytest, ytest_pred) print(mse)
mse = mean_squared_error(ytrain_pred,ytrain)
print(mse)
```

0.5289841670367244


0.5234413607125447

```
In [33]: mse = mean_squared_error(ytest, ytest_pred)
print(mse)
```

0.5289841670367244

In

```
[34]: plt.scatter(ytrain ,ytrain_pred,c='blue',marker='o',label='Training data')
plt.scatter(ytest,ytest_pred ,c='lightgreen',marker='s',label='Test data')
plt.xlabel('True values')
plt.ylabel('Predicted')
plt.title("True value vs Predicted value")
plt.legend(loc= 'upper left')
#plt.hlines(y=0,xmin=0,xmax=50)
plt.plot()
plt.show()
```



Name- Sharayu Nanaware Rollno:- 13243